

CSE6234 Software Design

Software Design Proposal and Software Design Pattern Document

HealthTech

Development of a Machine Learning-Based System for Cardiopulmonary Sound Separation

Group 1

Tutorial Section: TT5L

Name	Student ID
AL-SALOUL, ASHRAF ALI HUSSEIN	1221303805
AHMAD AKMAL ASYRAAF BIN SHAMS	242UC244CJ
RESHMA A/P KRISHNAMURTHY	243UC247KW
SHARWIN A/L R SIVALINGAM	241UC241C1

June 24, 2026

Contents

Abstract / Executive Summary	iv
1 Introduction	1
1.1 Executive Summary / Abstract	1
1.2 Background / Introduction	1
1.3 Problem Statement	2
1.4 Project Objectives	2
1.4.1 Literature Review	3
1.5 Project and System Scope	4
1.5.1 In Scope	4
1.5.2 Out of Scope	5
2 System Overview	6
2.1 Generic Use Case	7
2.1.1 Use Case Specifications	7
2.2 Software Architecture	13
2.3 Component Level Diagram	14
2.4 Class Diagram	16
2.5 Entity Relationship Diagram (ERD)	18
2.6 Object Diagram	19
3 Requirements	21

3.1	Functional Requirements	21
3.2	Non-Functional Requirements	22
3.3	Software Design Concepts	23
3.3.1	Abstraction	23
3.3.2	Modularity	24
3.3.3	Separation of Concerns	24
3.3.4	Information Hiding	24
3.3.5	Low Coupling and High Cohesion	25
3.4	The Five SOLID Principles	25
3.4.1	S – Single Responsibility Principle (SRP)	25
3.4.2	O – Open/Closed Principle (OCP)	26
3.4.3	L – Liskov Substitution Principle (LSP)	27
3.4.4	I – Interface Segregation Principle (ISP)	28
3.4.5	D – Dependency Inversion Principle (DIP)	28
3.4.6	Persistence Separation Between SQLite and Filesystem	29
3.5	Software Design Plan	30
4	Proposed Design Patterns	31
4.1	Facade Pattern	31
4.1.1	Pattern Type	31
4.1.2	Problem	31
4.1.3	Solution	32
4.1.4	Structure	32
4.1.5	Participants / Structure Explanation	33
4.1.6	Project Application	33
4.2	Strategy Pattern	36
4.2.1	Pattern Type	36

4.2.2	Problem	36
4.2.3	Solution	36
4.2.4	Structure	37
4.2.5	Participants / Structure Explanation	37
4.2.6	Project Application	38
4.3	Factory Method Pattern	40
4.3.1	Pattern Type	40
4.3.2	Problem	40
4.3.3	Solution	41
4.3.4	Structure	41
4.3.5	Participants / Structure Explanation	42
4.3.6	Project Application	42
5	Conclusion and Suggestions	45
5.1	Conclusion	45
5.2	Suggestions / Future Work	46
	Bibliography / References	47

Abstract / Executive Summary

This report presents a software design proposal and design pattern document for a HealthTech cardiopulmonary sound separation prototype. The proposed system accepts mixed WAV audio, uses NeoSSNet to separate the recording into heart sound and lung sound outputs, stores structured metadata and processing records in SQLite, and keeps uploaded and generated audio files in local storage.

The report focuses on software design rather than clinical validation. It documents the problem statement, objectives, system scope, UML models, requirements, design principles, and selected design patterns. FastAPI is used as the backend API layer, SQLite supports jobs, results, logs, and metadata, and the local filesystem stores audio artifacts and model files. The Facade, Strategy, and Factory Method patterns are proposed to simplify the separation workflow, isolate algorithm execution behind a strategy context, and separate upload validator creation from route logic.

Introduction

1.1 Executive Summary / Abstract

This report presents a software design proposal and design pattern document for a cardiopulmonary sound separation system. The proposed system is a HealthTech prototype that accepts mixed WAV audio and separates it into heart sound and lung sound outputs using NeoSS-Net. The project is designed as a local web-based application with a FastAPI backend, a simple HTML/CSS/JavaScript interface, SQLite database storage for metadata and processing records, and local filesystem storage for uploaded and generated audio files.

The purpose of the system is to demonstrate how machine learning inference can be integrated into a reusable software system rather than being kept only as an isolated model script. The application supports upload, validation, real separation, output preview, download, and processing history. This document focuses on software design quality, including requirements, UML modelling, database structure, layered architecture, and design patterns. The prototype is not a clinical diagnosis system and does not claim clinical validation. Instead, it is intended to show a practical and explainable software design for cardiopulmonary sound separation.

1.2 Background / Introduction

Cardiopulmonary sound recordings often contain both heart and lung components because the two sound sources are captured from nearby areas of the chest. In a mixed recording, the sound of one organ can interfere with analysis of the other. This creates a useful software engineering problem: the system must accept an input audio file, apply a sound separation model, produce separate heart and lung audio outputs, and keep the processing records traceable for review.

Recent research has shown that computational methods can support heart and lung sound processing, including source separation and deep learning-based heart sound analysis. However, a model alone is not enough for a usable application. A student, lecturer, researcher, or demonstration user needs a complete workflow that includes file upload, validation, inference, storage, result preview, download, error handling, and history. For this reason, the project is framed as

a software design assignment that connects machine learning with maintainable backend and frontend architecture.

The proposed system uses NeoSSNet as the core separation model and stores model-related files in local storage. SQLite is used to store structured data such as uploaded audio metadata, model records, separation jobs, result paths, metrics, and logs. The actual WAV audio files are stored on the filesystem so that the database remains lightweight and easy to inspect. This separation between database records and audio artifacts supports a practical local deployment suitable for a Final Year Project prototype.

1.3 Problem Statement

Mixed cardiopulmonary audio is difficult to use directly when the target task requires separate heart and lung sound outputs. If a system only stores uploaded mixed audio, users cannot easily preview the isolated components or use them for later analysis and demonstration. At the same time, if separation is implemented only as a standalone script, the workflow may lack upload handling, database tracking, result retrieval, logging, and a clear user interface.

The project therefore needs a reusable software system that can connect audio upload, NeoSSNet inference, file storage, SQLite metadata, and browser-based result viewing. The design must remain simple enough for a student project, but structured enough to explain through software design concepts, UML diagrams, and design patterns. The system must not present separated audio as clinical diagnosis; it should be described as a prototype that demonstrates automated cardiopulmonary sound separation and supporting software architecture.

1.4 Project Objectives

The objectives of this project are:

1. To design a local web-based system that accepts mixed cardiopulmonary WAV audio through a browser interface.
2. To validate uploaded audio files before they are stored or processed by the backend.
3. To integrate NeoSSNet as the real cardiopulmonary sound separation model for generating heart and lung output audio.
4. To save separated heart and lung WAV files in local output folders for preview and download.

5. To store structured metadata, job records, result paths, logs, and optional evaluation metrics in SQLite.
6. To provide result viewing, audio preview, download support, and processing history through the web interface.
7. To document a maintainable software design using UML diagrams, requirements, design principles, and selected design patterns.

1.4.1 Literature Review

Cardiopulmonary sound analysis studies information contained in heart and lung recordings, where useful physiological sounds may overlap because both sources are captured from the chest area. In practical recordings, the heart component can interfere with lung sound analysis, while lung sounds and breathing noise can also obscure heart sound patterns. Sound separation is therefore an important preprocessing and application problem because it attempts to recover clearer heart and lung signals from a mixed recording. For the proposed HealthTech prototype, the literature is relevant not only because it motivates machine learning-based separation, but also because it shows the need for a usable software system around the algorithm. A practical system must support upload, validation, inference, storage, preview, download, and processing history rather than treating separation as an isolated research script.

Recent studies show two complementary technical directions. The first is deep learning-based source separation. NeoSSNet has been proposed for real-time neonatal chest sound separation, with the aim of separating mixed chest sound recordings into heart and lung components (Poh et al., 2024). Its audio source separation approach is directly aligned with this project because the application must generate two usable WAV outputs: one heart sound file and one lung sound file. NeoSSNet also motivates the software architecture because its inference process can be placed behind a dedicated machine learning boundary and called from the backend service layer. However, the research contribution mainly concerns the model and separation performance. It does not by itself provide a complete application workflow for user upload, database-backed job tracking, browser preview, or maintainable backend organization.

The second direction is hybrid separation using signal processing and decomposition methods. A DAE–NMF–VMD method has been studied for separating heart and lung sounds by combining deep autoencoder-based representation learning, nonnegative matrix factorization, and variational mode decomposition (Sun et al., 2024). Compared with a neural separation model such as NeoSSNet, this approach shows that cardiopulmonary sound separation can also be designed as a multi-stage processing pipeline. This comparison is useful for the proposed system because future versions may need to compare algorithms with different preprocessing steps,

configuration parameters, model files, and output handling. Therefore, the software should avoid embedding one algorithm too deeply into route logic. A modular design with an interchangeable algorithm interface is better suited for future experimentation while keeping the current prototype focused on NeoSSNet.

Broader deep learning work in heart sound analysis further supports the need for careful system design. Deep learning techniques, datasets, and applications for heart sound analysis have been reviewed across preprocessing, feature extraction, model development, and potential clinical application areas (Zhao et al., 2024). This wider literature shows the growing importance of intelligent biomedical sound processing, but it also highlights concerns such as dataset limitations, generalisation, and the need for further refinement before broad clinical use. These points are important for this report because the proposed system should not be presented as a clinically validated diagnostic tool. It is more accurately described as a proof-of-concept HealthTech software prototype that demonstrates cardiopulmonary sound separation and the supporting software architecture needed to make the workflow usable and traceable.

Taken together, the studies show that cardiopulmonary sound separation and heart sound analysis have promising algorithmic foundations, but they also reveal a software design gap. Existing work focuses strongly on algorithms, while less attention is given to reusable application architecture for a local prototype. In particular, the literature does not fully address database-backed processing history, structured metadata, user upload workflow, separated audio preview and download, local file storage, logging, and maintainability through design patterns. The proposed system responds to this gap by combining FastAPI for the backend API, SQLite for metadata and job records, local filesystem storage for uploaded and separated WAV files, and NeoSSNet for real separation. This design connects research-level separation models with a practical software workflow that can be demonstrated, inspected, and extended in a student software design assignment.

1.5 Project and System Scope

1.5.1 In Scope

The project scope covers a local prototype that accepts mixed WAV audio, validates the file, stores the uploaded audio in a local folder, runs NeoSSNet-based separation, generates separated heart and lung WAV outputs, stores output files locally, records metadata and job status in SQLite, and provides browser-based preview, download, and processing history. The report also covers software design artefacts such as use case modelling, class modelling, database modelling, requirements, design principles, and design patterns.

The dataset and runtime storage are treated as separate concerns. Dataset files such as HLS-CMDS data are used for development or evaluation purposes, while runtime uploads and generated outputs are stored under the application storage folders. This separation avoids mixing training or testing data with user-uploaded demo files.

1.5.2 Out of Scope

The project does not aim to provide medical diagnosis, clinical decision support, hospital deployment, or validated clinical accuracy. It does not claim that separated audio is suitable for real clinical use without further testing and approval. Authentication, cloud deployment, multi-user role management, and production-scale monitoring are also outside the current scope unless they are added in a later phase.

Training a new model from runtime uploads is also outside the current system scope. Runtime uploads are used for inference and demonstration only. The current design focuses on integrating real NeoSSNet inference into a clean software system that is practical for local execution and clear enough to defend in a software design presentation.

System Overview

The proposed system is a local web-based prototype for cardiopulmonary sound separation. A user uploads a mixed WAV audio file through the browser interface, and the FastAPI backend validates the upload, stores the original file, creates database records, runs NeoSSNet separation, saves the separated heart and lung audio files, and returns results for preview and download. SQLite is used for structured records such as uploaded audio metadata, model information, separation jobs, result paths, metrics, and logs, while local folders are used for the actual audio files and model checkpoints.

The following diagrams describe the system from different software design viewpoints. The use case diagram presents the system functions from the user's perspective, the architecture diagram shows the main runtime layers, the component level diagram shows how modules interact, the class diagram shows the planned software components, the entity relationship diagram shows the database structure, and the object diagram gives an example of one completed separation job.

2.1 Generic Use Case




Figure 2.1: Generic Use Case Diagram

Figure 2.1 shows the main actions available to a single generic actor, the User. The diagram is intentionally simple because the same user-facing workflow can be demonstrated by a student, lecturer, examiner, or researcher without changing the system boundary. The central use cases are uploading mixed audio, running cardiopulmonary sound separation, previewing separated heart and lung outputs, downloading separated outputs, and viewing processing history.

The simplified diagram keeps internal responsibilities such as validation, model selection, metadata storage, and logging out of the actor view. Those responsibilities are still required by the backend, but they support the visible workflow rather than appearing as separate user goals. This makes the use case model easier to read while still showing how upload, separation, preview, download, and history are connected.

2.1.1 Use Case Specifications

A use case describes how an actor interacts with a system to achieve a specific goal. In software design, use cases are useful because they clarify system boundaries, user intentions, required inputs, normal workflows, and failure handling before detailed implementation begins. In the following use case specifications, the main course represents the normal successful flow, the alternate course represents another valid path through the same use case, and the exception courses describe error or failure conditions that the system should handle.

Table 2.1: RML-Style Use Case Specification for UC_001

Field	Specification
Name	Upload Mixed Audio
ID	UC_001
Description	The User uploads a mixed cardiopulmonary WAV file through the web-based interface so that it can be validated, stored, and prepared for separation.
Actor	User
Organization Benefits	Supports the first step of the prototype workflow by allowing mixed audio recordings to enter the system in a controlled and traceable way. It also creates metadata needed for later separation, result viewing, and history.
Frequency of Use	Occurs whenever the User wants to process a new mixed cardiopulmonary audio file during testing, demonstration, or review.
Triggers	The User selects a WAV file in the upload form and submits the upload request.
Preconditions	The web interface is available; the User has a mixed WAV audio file; local upload storage and SQLite are available.
Postconditions	A valid upload is saved in local storage, and a corresponding uploaded audio record is stored in SQLite with file name, path, timestamp, and metadata where available.
Main Course	1. The User opens the web interface. 2. The User selects a mixed WAV audio file. 3. The User submits the upload form. 4. The system validates the file type and audio properties. 5. The system stores the file in the upload folder. 6. The system records upload metadata in SQLite. 7. The system returns the uploaded audio identifier for separation.
Alternate Course	The User may upload another mixed WAV file after a successful upload. The system treats the second file as a new upload and creates a separate database record.
Exception Courses	If the file is missing, not a WAV file, unreadable, or unsupported, the system rejects the upload and shows a clear error message. If storage or database writing fails, the system does not continue to separation and records the failure where possible.

Table 2.2: RML-Style Use Case Specification for UC_002

Field	Specification
Name	Run Sound Separation
ID	UC_002
Description	The User starts cardiopulmonary sound separation for an uploaded mixed WAV file. The system uses the selected model if a model is chosen, or the active default model if no model is selected. NeoSSNet is the current real available separation model.
Actor	User
Organization Benefits	Demonstrates the main HealthTech function of the prototype by converting a mixed cardiopulmonary input into separate heart and lung sound outputs using real machine learning inference.
Frequency of Use	Occurs after a successful upload whenever the User wants to generate separated heart and lung outputs.
Triggers	The User clicks the separation action for an uploaded audio record and may choose an available model from the model dropdown.
Preconditions	A valid uploaded audio record exists; the uploaded WAV file is available in local storage; SQLite is available; the selected or active model record exists; the NeoSSNet checkpoint and configuration files are available.
Postconditions	A separation job record is created or updated, heart and lung WAV output files are stored locally, result metadata is saved in SQLite, and the job status shows completion or failure.
Main Course	1. The User selects or accepts the default separation model. 2. The User starts separation for an uploaded audio file. 3. The system retrieves the upload record and model record from SQLite. 4. The system creates a running separation job. 5. The system runs the selected or default separation model. 6. NeoSSNet generates heart and lung WAV output files. 7. The system stores output paths and job status in SQLite. 8. The system returns the completed job information to the web interface.
Alternate Course	If the User does not choose a model, the system uses the active default model from SQLite. In the current prototype, this default model is NeoSSNet.
Exception Courses	If the uploaded file is missing, the selected model does not exist, the checkpoint path is unavailable, or inference fails, the system marks the job as failed and stores an error message. The system does not create fake heart or lung outputs.

Table 2.3: RML-Style Use Case Specification for UC_003

Field	Specification
Name	Preview Separated Audio
ID	UC_003
Description	The User previews the separated heart and lung WAV outputs in the browser after a successful separation job.
Actor	User
Organization Benefits	Allows quick inspection of separation results during demonstration or testing without requiring external audio software. This improves usability and makes the prototype easier to evaluate.
Frequency of Use	Occurs after separation completes successfully and whenever the User reopens a previous result from history.
Triggers	The system displays a completed result, or the User opens a completed job from processing history.
Preconditions	A separation job has completed successfully; result records exist in SQLite; heart and lung output files exist in local storage; the browser supports playback of the generated WAV files.
Postconditions	The User can listen to the separated heart and lung outputs through browser audio controls. No database record is changed by playback alone.
Main Course	1. The User opens the result view for a completed job. 2. The system loads the stored heart and lung output paths from SQLite. 3. The system displays browser audio controls for the separated outputs. 4. The User plays the heart sound output. 5. The User plays the lung sound output. 6. The User compares the separated outputs with the original mixed audio if the original preview is available.
Alternate Course	The User may preview only one output, such as the heart sound, and still proceed to download or history without playing the other output.
Exception Courses	If an output file is missing, the system should avoid showing a broken playback experience and should display a clear unavailable result state. If the browser cannot play the file, the download option may still be used for external playback.

Table 2.4: RML-Style Use Case Specification for UC_004

Field	Specification
Name	Download Separated Audio
ID	UC_004
Description	The User downloads the generated heart or lung WAV output file from a completed separation job.
Actor	User
Organization Benefits	Supports offline review, presentation, comparison, and further analysis of generated separation outputs while keeping the runtime application simple and local.
Frequency of Use	Occurs whenever the User wants to keep or share a separated output after previewing or viewing a completed job.
Triggers	The User selects the download action for the heart output or lung output.
Preconditions	A completed separation result exists; the selected output file path is stored in SQLite; the requested heart or lung WAV file exists in local output storage.
Postconditions	The selected output file is delivered to the browser as a downloadable WAV file. The stored result record remains unchanged.
Main Course	1. The User opens a completed job result. 2. The User selects either the heart output download or lung output download. 3. The system checks the related result record. 4. The system locates the requested WAV file in local storage. 5. The system sends the file to the browser for download. 6. The User saves or opens the downloaded file.
Alternate Course	The User may download both heart and lung outputs from the same completed job. Each download is handled as a separate file request.
Exception Courses	If the job ID is invalid, the output type is not recognised, or the requested file is missing from local storage, the system returns a clear error response instead of an empty or incorrect file.

Table 2.5: RML-Style Use Case Specification for UC_005

Field	Specification
Name	View Processing History
ID	UC_005
Description	The User views saved processing history from SQLite, including previous uploads, separation jobs, model information, statuses, result links, and timestamps.
Actor	User
Organization Benefits	Provides traceability for the prototype by showing how uploaded files, selected models, generated outputs, and job statuses are connected. This supports demonstration, debugging, and report explanation.
Frequency of Use	Occurs whenever the User wants to review previous processing activity or reopen a completed result.
Triggers	The User opens the history section or refreshes the main web interface history view.
Preconditions	The FastAPI backend and SQLite database are available. History may contain previous records, or it may be empty in a fresh setup.
Postconditions	The web interface displays available processing history records, including completed, failed, or pending jobs where applicable.
Main Course	1. The User opens the processing history section. 2. The system queries saved upload, job, model, and result records from SQLite. 3. The system formats the records for the web interface. 4. The system displays job status, original file information, model used, timestamps, and result actions where available. 5. The User reviews previous activity or opens a completed result.
Alternate Course	If no previous jobs exist, the system displays an empty history state and keeps the upload workflow available.
Exception Courses	If the database cannot be read or a history record points to a missing output file, the system should display the available metadata with a clear error state for the affected result instead of failing the whole page.

2.2 Software Architecture




Figure 2.2: Software Architecture Diagram

Figure 2.2 shows the overall architecture of the prototype. The user interacts with a browser-based web interface, which sends RESTful HTTP requests to the FastAPI application. The FastAPI application is separated into route handling, business services, and the machine learning boundary. SQLite stores structured metadata such as uploads, jobs, result paths, metrics, and logs, while local storage folders keep uploaded WAV files, separated heart and lung outputs, model files, and runtime logs.

2.3 Component Level Diagram




Figure 2.3: Component Level Diagram

Figure 2.3 shows the main software components and their responsibilities. The UI layer handles upload, preview, download, and history screens. The FastAPI router layer receives requests for upload, separation, model listing, result retrieval, download, and history. The business service layer coordinates validation, separation workflow, model lookup, file storage, and result persistence. The ML layer runs separation through `SeparationEngine`, `SeparationAlgorithm`, and `NeoSSNetStrategy`. SQLite stores processing records, while local storage keeps the actual audio and model files.

2.4 Class Diagram




Figure 2.4: Overall System UML Class Diagram

Figure 2.4 presents the main classes and interfaces in the cardiopulmonary sound separation system. Each class or interface appears only once, and the diagram shows the API router modules, facade services, strategy classes, upload validation factory classes, and SQLite data model classes. The diagram uses UML class-style boxes, selected attributes or methods, associations, dependencies, and interface implementation arrows to describe the whole system structure.

Design pattern roles are separated so that one main class is not overloaded with multiple pattern labels. `SeparationService` is shown as the Facade used by `SeparationRouter`. `SeparationEngine` is shown as the Strategy Context that depends on `SeparationAlgorithm`, and `NeoSSNetStrategy` is the current Concrete Strategy. `UploadRouter` is shown as the Factory Method client, `AudioValidatorFactory` is the Creator, `WavAudioValidatorFactory` is the ConcreteCreator, `AudioValidator` is the Product interface, and `WavAudioValidator` is the ConcreteProduct.

2.5 Entity Relationship Diagram (ERD)




Figure 2.5: Entity Relationship Diagram

Figure 2.5 shows the SQLite database structure used to store structured project data. The `uploaded_audio` table records each uploaded WAV file, the `model` table records available

model information, and the `separation_job` table connects an uploaded file with the model used for processing. The `separation_result` table stores the paths and metadata for generated heart and lung outputs.

The design intentionally stores file paths and metadata in SQLite rather than storing large WAV audio blobs in the database. This supports a practical student project architecture: local storage keeps uploaded and generated audio files, while SQLite supports history, result lookup, job status, metrics, and logs. As a result, the browser can display previous jobs and download links without scanning folders manually.

2.6 Object Diagram




Figure 2.6: Object Diagram for a Completed Separation Job

Figure 2.6 gives a concrete example of how the system data may look after one successful separation. One `UploadedAudio` object is linked to one `SeparationJob`, the job uses an active `Model`, and the completed job produces a `SeparationResult` containing heart and lung

output file paths.

The object diagram helps explain how upload, separation, storage, result viewing, and history are connected at runtime. When a user views history, the system can retrieve the job, original upload metadata, model information, output paths, optional metrics, and log record as related objects. This gives the project a traceable flow from the original mixed input to the separated outputs.

Requirements

This chapter defines the main functional requirements, non-functional requirements, software design concepts, and design principles for the cardiopulmonary sound separation system. The system is treated as a prototype and proof of concept for a software design assignment and Final Year Project context. It is designed to demonstrate a realistic local web application that can upload mixed cardiopulmonary audio, run NeoSSNet-based separation, store metadata, and present results clearly to students, lecturers, researchers, or demonstration users. The requirements do not claim clinical validation; they focus on software behaviour, maintainability, and demonstration readiness.

3.1 Functional Requirements

Functional requirements describe the main behaviours that the system must provide to the user and to the backend workflow.

Table 3.1: Functional Requirements

ID	Requirement	Description
FR-01	Upload mixed WAV audio	The user shall be able to select and upload a mixed cardiopulmonary WAV file through the web interface.
FR-02	Validate audio file	The backend shall check that the uploaded file is a valid supported audio file before saving it for processing.
FR-03	Run cardiopulmonary sound separation	The backend shall run the separation workflow using NeoSSNet to separate the mixed input into heart and lung sound outputs.
FR-04	Generate heart and lung output files	The system shall save separated heart and lung audio files in the configured output folders.
FR-05	Preview separated audio	The web interface shall allow the user to preview the original audio and the separated heart and lung outputs where browser playback is supported.
FR-06	Download separated audio	The user shall be able to download the generated heart and lung output files for later review or demonstration.
FR-07	View processing history	The system shall show previous uploads, jobs, statuses, and result links so that users can review earlier processing activity.
FR-08	Store logs and metadata	The system shall store upload metadata, job records, result paths, metrics where available, and important processing logs in SQLite.

These functional requirements support the main demonstration flow of the project. A user uploads a mixed sound, the backend validates and processes it, the system stores the generated files, and the browser presents previews, downloads, and history. This keeps the project focused on a working end-to-end software system rather than only a standalone machine learning script.

3.2 Non-Functional Requirements

Non-functional requirements describe the quality attributes expected from the prototype. These qualities are important because the system should be understandable, maintainable, and reliable enough for local testing and presentation.

Table 3.2: Non-Functional Requirements

Quality Attribute	Requirement
Usability	The interface should make the upload, separation, preview, download, and history flow clear for a student or examiner during a demonstration.
Maintainability	The code should be organised into clear frontend, router, service, machine learning, database, and storage responsibilities so that future changes are easier to make.
Reliability	The backend should handle invalid files, failed separation jobs, missing outputs, and database errors with clear status messages and logs.
Modularity	Each major part of the system should be replaceable or testable with minimal impact on unrelated components.
Performance	The system should process short demonstration WAV files within a reasonable time on a local machine, while avoiding unnecessary database storage of large audio blobs.
Data integrity	SQLite records should consistently link uploaded audio, model records, separation jobs, results, metrics, and logs. Failed jobs should keep useful error information.
Portability and local execution	The system should run locally using Python, FastAPI, SQLite, PyTorch, and local folders without requiring cloud services or complex deployment infrastructure.

The non-functional requirements guide implementation decisions. For example, local execution supports a practical student project setup, while data integrity and logging support traceability when explaining how a mixed input becomes separated heart and lung outputs.

3.3 Software Design Concepts

3.3.1 Abstraction

Abstraction is used to hide detailed processing steps behind simpler operations. In this project, the separation route does not need to know how model records are selected, how NeoSSNet is called, or how heart and lung output paths are written. It can call a service-level operation such as `separate_uploaded_audio` and treat the workflow as one meaningful software action.

The machine learning boundary also uses abstraction. `SeparationAlgorithm` represents the expected behaviour of a separation model, while `NeoSSNetStrategy` provides the current real implementation. This allows the service layer to express the workflow in terms of cardiopulmonary sound separation rather than in terms of one concrete model class.

3.3.2 Modularity

Modularity divides the system into smaller parts with clear responsibilities. The frontend handles upload controls, audio preview, download links, and history display. FastAPI routers handle HTTP requests. Services coordinate workflows. The ML layer performs real separation. SQLite models store structured records, and local folders store uploaded and generated WAV files.

This modular structure is useful for the prototype because a change in one area does not require rewriting the entire system. For example, improving the frontend model dropdown should not require changing NeoSSNet inference, and changing output folder naming should not require changing the database schema.

3.3.3 Separation of Concerns

Separation of concerns prevents unrelated responsibilities from being mixed together. The upload route receives files, the storage service manages local file placement, the model service reads available model records, the separation service coordinates processing, and the ML strategy runs inference. This keeps web request handling, file management, model selection, and audio separation as separate concerns.

The same idea applies to data storage. Runtime upload folders, generated output folders, model checkpoints, dataset folders, and SQLite metadata serve different purposes and remain separate. This design avoids training from runtime upload folders and avoids storing large WAV content directly in SQLite.

3.3.4 Information Hiding

Information hiding means that internal implementation details are kept inside the module or class responsible for them. File validation rules should stay inside the upload validation component, storage paths should stay inside `StorageService`, model lookup should stay inside `ModelService`, and NeoSSNet inference details should stay inside `NeoSSNetStrategy` and the existing NeoSSNet inference function. `SeparationEngine` only sees the shared `SeparationAlgorithm` interface.

This protects the rest of the system from unnecessary implementation knowledge. For example, the route does not need to know where the checkpoint file is stored or how the strategy calls the underlying inference function; it only needs the job result returned by the service.

3.3.5 Low Coupling and High Cohesion

Low coupling means that components should depend on each other only where necessary. In this project, SeparationEngine depends on the SeparationAlgorithm abstraction rather than depending directly on NeoSSNet-specific workflow logic. This reduces the impact of future model changes because the upload, result, history, and download routes can remain stable.

High cohesion means that each component should focus on related tasks. StorageService is cohesive when it handles file storage concerns, ModelService is cohesive when it handles model lookup, and NeoSSNetStrategy is cohesive when it only adapts NeoSSNet inference to the shared algorithm interface. Cohesion helps the system remain easy to explain and test.

3.4 The Five SOLID Principles

3.4.1 S – Single Responsibility Principle (SRP)

A class should have only one reason to change, meaning it should have a single, well-defined job. In this project, StorageService handles storage-related work such as file paths, saving uploaded WAV files, and building output paths. It should not run NeoSSNet inference or decide which model is active.

Listing 3.1: SRP Example: Storage Responsibility Only

```
class StorageService:
    # SRP: This class has one responsibility: file storage.
    # This class should change only when file storage behavior changes.
    def save_uploaded_wav(self, upload_file) -> str:
        # This method belongs to the same storage responsibility.
        stored_path = self.build_upload_path(upload_file.filename)
        self.write_file(upload_file, stored_path)
        return stored_path

class SeparationService:
    # SRP: Workflow coordination stays outside StorageService.
    def upload_then_record(self, file, db):
        stored_path = StorageService().save_uploaded_wav(file)
        return create_uploaded_audio_record(db, stored_path)
```

The single responsibility is storage. What is separated is the database workflow and ML inference. This means a change to the upload folder naming affects `StorageService`, while a change to `NeoSSNet` affects `NeoSSNetStrategy` instead.

3.4.2 O – Open/Closed Principle (OCP)

Software entities should be open for extension but closed for modification, allowing behavior to change without altering existing code. In this project, a new implemented separation algorithm should be added as another `SeparationAlgorithm` implementation without rewriting `SeparationEngine`.

Listing 3.2: OCP Example: Extend with a New Algorithm

```
class SeparationAlgorithm(Protocol):
    # OCP: This interface is open for extension.
    def separate(self, **paths):
        ...

class SeparationEngine:
    # OCP: This class is closed for modification.
    def __init__(self, algorithm: SeparationAlgorithm):
        self.algorithm = algorithm

    def separate(self, paths):
        return self.algorithm.separate(
            input_wav_path=paths.input_wav,
            model_path=paths.model_path,
            model_config_path=paths.config_path,
            heart_output_path=paths.heart_output,
            lung_output_path=paths.lung_output,
        )

class NeoSSNetStrategy:
    # OCP: This new class extends behavior without modifying
    # SeparationEngine.
    def separate(self, input_wav_path, model_path, model_config_path,
                  heart_output_path, lung_output_path):
        return run_neossnet_inference(...)
```

The open part is the set of concrete SeparationAlgorithm implementations. The closed part is SeparationEngine; it can keep the same call structure while a future real algorithm is added through a new strategy class and service-level selection logic.

3.4.3 L – Liskov Substitution Principle (LSP)

Subtypes must be substitutable for their base types without altering the correctness of the program. NeoSSNetStrategy can replace any object typed as SeparationAlgorithm because it provides the expected separate(...) method and returns the expected result object.

Listing 3.3: LSP Example: NeoSSNet Substitutes the Strategy Interface

```
class SeparationAlgorithm(Protocol):
    # LSP: This is the parent/base abstraction.
    def separate(self, input_path: str):
        ...

class NeoSSNetStrategy:
    # LSP: This is the child/subtype implementation.
    def separate(self, input_path: str):
        return run_neossnet_inference(input_path)

class SeparationEngine:
    # LSP: SeparationEngine accepts the parent/base type.
    def __init__(self, algorithm: SeparationAlgorithm):
        self.algorithm = algorithm

    def run(self, input_path: str):
        return self.algorithm.separate(input_path)

# LSP: The child can replace the parent safely.
engine = SeparationEngine(NeoSSNetStrategy())
```

The substitutable type is NeoSSNetStrategy. The base abstraction is SeparationAlgorithm. Correctness is preserved because SeparationEngine still receives separated heart and lung output metadata in the same result format.

3.4.4 I – Interface Segregation Principle (ISP)

Clients should not be forced to depend on methods they do not use, favoring small, specific interfaces over large, general ones. The project uses small interfaces such as `AudioValidator` for upload validation and `SeparationAlgorithm` for separation inference.

Listing 3.4: ISP Example: Small Focused Interfaces

```
class AudioValidator(Protocol):
    # ISP: This is a small focused interface.
    async def validate(self, upload_file) -> None:
        ...

class SeparationAlgorithm(Protocol):
    # ISP: This is another small focused interface.
    def separate(self, input_wav_path, model_path, model_config_path,
                  heart_output_path, lung_output_path):
        ...

class UploadRouter:
    # ISP: UploadRouter only depends on validate().
    async def upload_audio(self, file, validator: AudioValidator):
        # ISP: The client is not forced to depend on unused methods.
        await validator.validate(file)
        return await StorageService.save_uploaded_wav(file)
```

The small interfaces are `AudioValidator` and `SeparationAlgorithm`. Upload routing is not forced to know about model inference, and the strategy engine is not forced to know about upload headers or database records.

3.4.5 D – Dependency Inversion Principle (DIP)

Depend on abstractions/interfaces rather than concrete implementations. In this project, `SeparationEngine` depends on `SeparationAlgorithm`, not directly on `NeoSSNetStrategy`. The concrete strategy is supplied after model selection is resolved.

Listing 3.5: DIP Example: Engine Depends on the Abstraction

```
class SeparationAlgorithm(Protocol):
```

```

# DIP: This is the abstraction.
def separate(self, **paths):
    ...

class NeoSSNetStrategy:
    # DIP: This is the low-level concrete implementation.
    def separate(self, **paths):
        return run_neosnet_inference(**paths)

class SeparationEngine:
    # DIP: This is the high-level module.
    def __init__(self, algorithm: SeparationAlgorithm):
        # DIP: The high-level module depends on the abstraction,
        # not the concrete class.
        self.algorithm = algorithm

    def separate(self, **paths):
        return self.algorithm.separate(**paths)

algorithm = selected_algorithm
engine = SeparationEngine(algorithm)

```

The abstraction depended on is `SeparationAlgorithm`. The concrete implementation is `NeoSSNetStrategy`. This keeps the high-level engine independent from NeoSSNet internals while still allowing the current real NeoSSNet implementation to run.

3.4.6 Persistence Separation Between SQLite and Filesystem

The design separates structured metadata from large binary audio files. SQLite stores records such as filenames, file paths, upload timestamps, job status, model information, output paths, metrics, and logs. The filesystem stores uploaded WAV files, generated heart and lung outputs, model checkpoints, and dataset files. This approach is simpler to run locally and avoids making the database unnecessarily large.

3.5 Software Design Plan

The proposed software design is organised into layers. This layered plan supports a clear implementation path and makes the system easier to explain in the report and viva presentation.

Table 3.3: Software Design Plan by Layer

Layer	Main Responsibility	Project Role
Frontend layer	HTML, CSS, JavaScript, templates, and browser audio controls	Provides upload form, status messages, original and separated audio playback, download links, and history display.
Backend API layer	FastAPI application and routers	Receives upload, separation, result, download, history, log, and health requests.
Service layer	Workflow coordination and business logic	Connects validation, database updates, NeoSSNet inference, output storage, logging, and result retrieval.
ML inference layer	Audio preprocessing and model inference	Loads or prepares audio, applies the NeoSSNet model, and returns separated heart and lung waveforms.
Database layer	SQLite schema and repository-style access	Stores uploaded audio metadata, model records, job states, result paths, metrics, and system logs.
Storage layer	Local filesystem folders	Stores uploaded mixed WAV files, temporary files, separated heart and lung outputs, logs where needed, and ML model weights.

The design plan follows the project priority of real functionality first, then simplicity and local execution. It avoids unnecessary cloud, microservice, or authentication complexity unless those features are explicitly added later.

Proposed Design Patterns

This chapter documents three design patterns used in the cardiopulmonary sound separation system. The roles are intentionally separated so that each pattern has its own purpose and its own main classes. `SeparationService` is explained as the Facade. `SeparationEngine` is explained as the Strategy Context. `AudioValidatorFactory` and `WavAudioValidatorFactory` are explained as the Factory Method Creator and `ConcreteCreator` used during upload validation.

4.1 Facade Pattern

The Facade Pattern provides a simplified interface to a group of related subsystem operations (Gamma et al., 1994). In this project, it is used to keep the FastAPI separation route small while the service layer coordinates model lookup, storage paths, result recording, and job status updates.

4.1.1 Pattern Type

The Facade Pattern is a structural design pattern. It organises a complex workflow behind one higher-level interface that is easier for a client to call.

4.1.2 Problem

The separation endpoint should not directly handle model lookup, upload lookup, output path preparation, job creation, inference coordination, result persistence, and error handling. Placing all of that logic inside `SeparationRouter` would make the route difficult to test and difficult to explain in the software design report.

4.1.3 Solution

SeparationRouter calls SeparationService through one high-level operation. SeparationService acts as the Facade and coordinates the service-level subsystems: ModelService, StorageService, and ResultService. The router remains responsible for HTTP request and response handling, while the Facade owns the separation workflow.

4.1.4 Structure




Figure 4.1: Facade Pattern Template Structure

1. **Client:** The caller that wants to use subsystem functionality without controlling every subsystem directly.
2. **Facade:** The simplified interface that coordinates subsystem work.
3. **Subsystems:** The internal components that perform specialised tasks.

4.1.5 Participants / Structure Explanation

Table 4.1: Facade Pattern Mapping

Pattern Participant	Project Class	Role in This Project
Client	SeparationRouter	Receives the <code>/separate/{audio_id}</code> request and calls the Facade.
Facade	SeparationService	Coordinates the high-level separation workflow and hides subsystem details from the route.
Subsystem	ModelService	Retrieves the selected model or active default model from SQLite.
Subsystem	StorageService	Resolves project paths and prepares heart/lung output paths.
Subsystem	ResultService	Creates and reads result records used by result, history, and download views.

4.1.6 Project Application

i. Function or Software Component Affected

The affected component is the backend separation workflow in `app/services/separation_service.py`. `SeparationRouter` in `app/routers/separation.py` delegates the workflow to the Facade instead of coordinating subsystem details itself.

ii. Description of Workflow or Data Flow

The user starts separation from the web interface. `SeparationRouter` receives the uploaded audio identifier and optional model identifier, then calls `SeparationService`. The Facade retrieves the upload and model, prepares output paths, creates a running job, coordinates the selected real separation strategy through the engine, records the result, and returns a response containing the generated heart and lung output paths. The Facade uses the strategy and model-selection helper internally, but its pattern role remains Facade only.

iii. Sample Class Diagram




Figure 4.2: Facade Pattern Project Class Diagram

iv. Facade Pattern Sequence Diagram




Figure 4.3: Facade Pattern Sequence Diagram

v. Sample Potential Code

Listing 4.1: Facade Pattern Code Example

```

class SeparationRouter:
    # Client: SeparationRouter calls the Facade.
    def separate_audio(self, audio_id: int, model_id: int | None, db):
        return SeparationService().separate_uploaded_audio(
  
```

```

        db=db,
        audio_id=audio_id,
        model_id=model_id,
    )

class SeparationService:
    # Facade: SeparationService provides one simple interface.
    def separate_uploaded_audio(self, db, audio_id, model_id=None):
        uploaded_audio = get_uploaded_audio(db, audio_id)
        # Subsystem: ModelService handles model selection.
        model = ModelService.get_model_for_separation(db, model_id)
        job = self.create_running_job(
            db, uploaded_audio.uploaded_audio_id, model.model_id
        )
        # Subsystem: StorageService handles file operations.
        output_paths = StorageService.build_separation_output_paths(job.job_id
        )
        output = run_selected_algorithm(uploaded_audio, model, output_paths)
        # Subsystem: ResultService handles result/history operations.
        result = ResultService.create_separation_result(db, job.job_id, output
        )
        return result

    def create_running_job(self, db, uploaded_audio_id, model_id):
        ...

    def mark_job_failed(self, db, job, error, processing_time_ms):
        ...

```

vi. Benefits

The Facade keeps route logic thin, makes the main separation use case easier to test, and gives the report a clear explanation of where workflow coordination happens. It also improves maintainability because storage, model lookup, and result persistence remain separate subsystem concerns.

vii. Limitations

The Facade can become too large if detailed validation, model internals, or result formatting are placed inside it. The implementation avoids this by keeping upload validation in its own component, model lookup in `ModelService`, output path handling in `StorageService`, and result creation in `ResultService`.

4.2 Strategy Pattern

The Strategy Pattern allows a family of algorithms to be used through a common interface (Gamma et al., 1994). In this project, it keeps the separation engine independent from NeoSSNet internals while still using the current real NeoSSNet inference path.

4.2.1 Pattern Type

The Strategy Pattern is a behavioral design pattern. It defines a common interface for interchangeable algorithms and lets a context call the selected algorithm through that interface.

4.2.2 Problem

The backend should not spread NeoSSNet-specific inference calls across routers or service workflow code. If the workflow directly calls `run_neosnet_inference`, the system becomes harder to test and harder to extend when another real implemented separation algorithm is added.

4.2.3 Solution

`SeparationEngine` is introduced as the Strategy Context. It depends on the `SeparationAlgorithm` abstraction and calls `algorithm.separate(...)`. `NeoSSNetStrategy` implements the abstraction and wraps the existing real `run_neosnet_inference` function without duplicating NeoSSNet code.

4.2.4 Structure




Figure 4.4: Strategy Pattern Template Structure

1. **Context:** The object that uses a strategy through a shared interface.
2. **Strategy:** The abstraction that defines the algorithm operation.
3. **ConcreteStrategy:** The implemented algorithm used at runtime.

4.2.5 Participants / Structure Explanation

Table 4.2: Strategy Pattern Mapping

Pattern Participant	Project Class	Role in This Project
Context	SeparationEngine	Runs the selected algorithm through the abstraction.
Strategy	SeparationAlgorithm	Defines the separate(...) contract and result shape.
ConcreteStrategy	NeoSSNetStrategy	Calls the real NeoSSNet inference boundary and returns heart/lung output metadata.

4.2.6 Project Application

i. Function or Software Component Affected

The affected component is the ML inference boundary under `app/ml/`. The relevant files are `separation_engine.py`, `separation_algorithm.py`, `neossnet_strategy.py`, and `neossnet_inference.py`.

ii. Description of Workflow or Data Flow

`SeparationService` obtains the selected model and prepares the algorithm object used for the current separation request. The service passes the algorithm to `SeparationEngine`. The engine calls `algorithm.separate(...)` with the uploaded WAV path, model checkpoint path, model configuration path, and heart/lung output paths. `NeoSSNetStrategy` then calls the existing `NeoSSNet` inference function and returns output metadata.

iii. Sample Class Diagram




Figure 4.5: Strategy Pattern Project Class Diagram

iv. Strategy Pattern Sequence Diagram




Figure 4.6: Strategy Pattern Sequence Diagram

v. Sample Potential Code

Listing 4.2: Strategy Pattern Code Example

```

class SeparationAlgorithm(Protocol):
    # Strategy: SeparationAlgorithm defines the common interface.
    def separate(self, input_wav_path, model_path, model_config_path,
                  heart_output_path, lung_output_path, device_name="cpu"):
        ...

class NeoSSNetStrategy:
    # ConcreteStrategy: NeoSSNetStrategy implements the algorithm.
    def separate(self, input_wav_path, model_path, model_config_path,
                  heart_output_path, lung_output_path, device_name="cpu"):
        return run_neossnet_inference(
            input_wav_path=input_wav_path,
            model_path=model_path,
            model_config_path=model_config_path,
            heart_output_path=heart_output_path,
            lung_output_path=lung_output_path,
            device_name=device_name,
        )

class SeparationEngine:

```

```
# Context: SeparationEngine uses a Strategy.
def __init__(self, algorithm: SeparationAlgorithm):
    self.algorithm = algorithm

def separate(self, **paths):
    return self.algorithm.separate(**paths)
```

vi. Benefits

The Strategy Pattern keeps NeoSSNet code isolated, makes the engine easy to test with a mocked algorithm, and allows the service workflow to remain stable when another real separation implementation is added later.

vii. Limitations

The current prototype mainly runs NeoSSNet, so the strategy layer must remain small. Over-complicating the abstraction would make the student project harder to maintain without improving the current working inference path.

4.3 Factory Method Pattern

The Factory Method Pattern creates products through a creator class instead of forcing the client to instantiate concrete products directly (Gamma et al., 1994). In this project, the pattern is used in the upload workflow to create the correct audio validator before the uploaded WAV file is saved.

4.3.1 Pattern Type

The Factory Method Pattern is a creational design pattern. It separates object creation from business workflow logic.

4.3.2 Problem

UploadRouter should not directly instantiate a concrete validator such as WavAudioValidator. If the route creates validators directly, upload request handling

becomes tightly coupled to one validation implementation. Adding another validator later would require changing the route instead of extending the validation creation logic.

4.3.3 Solution

AudioValidatorFactory declares the factory method `create_validator(filename)` and returns an `AudioValidator` product. WavAudioValidatorFactory is the concrete creator used by the current prototype. When the uploaded filename ends with `.wav`, it creates and returns a `WavAudioValidator`. UploadRouter uses the returned object through the `AudioValidator` product interface.

4.3.4 Structure




Figure 4.7: Factory Method Pattern Template Structure

1. **Client:** The code that needs a product but should not create the concrete product directly.
2. **Creator:** The class that declares the factory method.
3. **ConcreteCreator:** The implementation that creates a concrete product.
4. **Product:** The interface returned to the client.
5. **ConcreteProduct:** The implementation created by the factory method.

4.3.5 Participants / Structure Explanation

Table 4.3: Factory Method Pattern Mapping

Pattern Participant	Project Class	Role in This Project
Client	UploadRouter	Receives the uploaded file and asks the factory for a validator.
Creator	AudioValidatorFactory	Declares <code>create_validator(filename)</code> and returns an <code>AudioValidator</code> .
ConcreteCreator	WavAudioValidatorFactory	Creates a WAV validator when the uploaded file extension is supported.
Product	AudioValidator	Defines the common <code>validate(upload_file)</code> operation used by the route.
ConcreteProduct	WavAudioValidator	Validates the <code>.wav</code> extension and RIFF/WAVE file header.

4.3.6 Project Application

i. Function or Software Component Affected

The affected component is upload validation in `app/routers/upload.py` and `app/services/audio_validation.py`. The factory creates a validator before the uploaded file is saved to local storage or inserted into SQLite.

ii. Description of Workflow or Data Flow

The user selects a WAV file and starts the upload flow. `UploadRouter` receives the file and calls `audio_validator_factory.create_validator(file.filename)`. The concrete creator, `WavAudioValidatorFactory`, checks the file extension and returns a `WavAudioValidator`. The route then calls `validator.validate(file)` through the `AudioValidator` product interface before saving the file.

iii. Sample Class Diagram




Figure 4.8: Factory Method Pattern Project Class Diagram

iv. Factory Method Pattern Sequence Diagram




Figure 4.9: Factory Method Pattern Sequence Diagram

v. Sample Potential Code

Listing 4.3: Factory Method Pattern Code Example

```

class AudioValidator(Protocol):
    # Product: common validation interface.
    async def validate(self, upload_file):
        ...
  
```

```

class WavAudioValidator:
    # ConcreteProduct: validates WAV uploads.
    async def validate(self, upload_file):
        ...

class AudioValidatorFactory(ABC):
    # Creator: declares the factory method.
    @abstractmethod
    def create_validator(self, filename):
        ...

class WavAudioValidatorFactory(AudioValidatorFactory):
    # ConcreteCreator: creates the WAV validator.
    def create_validator(self, filename):
        if filename.endswith(".wav"):
            return WavAudioValidator()
        raise ValueError("Only .wav audio files are accepted.")

class UploadRouter:
    # Client: asks the Creator for a Product.
    async def upload_audio(self, file):
        validator = audio_validator_factory.create_validator(file.filename)
        await validator.validate(file)

```

vi. Benefits

The Factory Method keeps validator creation separate from upload request handling. It makes the upload route easier to maintain because the route depends on the `AudioValidator` product interface rather than directly constructing `WavAudioValidator`.

vii. Limitations

The current prototype has one concrete validator, `WavAudioValidator`, because the first version accepts only WAV files. The factory should remain small unless more audio formats or validation rules are added later.

Conclusion and Suggestions

5.1 Conclusion

This report proposed the design of a HealthTech cardiopulmonary sound separation system as a prototype and proof of concept. The system is intended to accept mixed WAV audio, run cardiopulmonary sound separation, and produce separate heart sound and lung sound output files. It supports the main user workflow of upload, validation, separation, preview, download, and processing history while keeping the implementation practical for a student software design project.

The proposed architecture is important because the machine learning model is only one part of the complete system. FastAPI provides the backend API layer for handling upload, separation, result, download, and history requests. NeoSSNet provides the core separation model for generating heart and lung audio outputs. SQLite stores structured metadata such as uploaded audio records, model details, separation jobs, result paths, evaluation metrics, and logs. Local storage keeps the actual uploaded WAV files, separated output files, and model weights, which avoids storing large audio content directly inside the database.

The selected design patterns strengthen the proposed design. The Facade Pattern simplifies the separation workflow by allowing the separation router to call a high-level service instead of directly managing model lookup, storage paths, result records, and database updates. The Strategy Pattern keeps algorithm execution inside SeparationEngine, which depends on the SeparationAlgorithm abstraction while NeoSSNetStrategy wraps the real NeoSSNet inference path. The Factory Method Pattern is used in upload validation, where UploadRouter asks AudioValidatorFactory for an AudioValidator product and the current concrete creator returns WavAudioValidator.

Overall, the proposed design demonstrates how machine learning, database persistence, local file storage, and web interaction can be organised into a reusable software system. The report does not claim clinical validation or real hospital deployment. Instead, it presents a clear and maintainable software design suitable for assignment submission, demonstration, and future implementation improvement.

5.2 Suggestions / Future Work

Several improvements can be considered after the core prototype is working:

1. **Improve frontend usability:** The user interface can be refined with clearer progress indicators, better error messages, improved history filtering, and more polished audio playback controls.
2. **Add more datasets and noise conditions:** Future testing can include broader cardiopulmonary recordings, different noise levels, and additional real-world recording conditions to better understand model behaviour.
3. **Add asynchronous or background processing:** Longer audio files or slower inference tasks can be handled using background jobs so that the web request does not need to wait synchronously for the full separation process.
4. **Add authentication if needed:** If the system is extended beyond a local demonstration, user login and role-based access can be added to protect uploaded files and processing history.
5. **Add algorithms through Strategy:** Alternative approaches, baseline methods, or future ONNX-based versions can be added as new strategies without changing the main router workflow.
6. **Improve evaluation metrics and reporting:** The system can provide clearer metric reports such as SNR, SDR, SI-SDR, MSE, or MAE where suitable reference data is available.
7. **Validate on broader real-world recordings:** Further validation should be performed on more varied recordings before making any claim about robustness or clinical usefulness.

In conclusion, the project provides a strong foundation for a practical software design around cardiopulmonary sound separation. Its main value is not only the use of NeoSSNet, but also the structured design that makes upload handling, inference, storage, logging, result viewing, and future extension easier to understand and maintain.

Bibliography / References

- Gamma, E., Helm, R., Johnson, R., & Vlissides, J. (1994). *Design patterns: Elements of reusable object-oriented software*. Addison-Wesley.
- Poh, Y. Y., Grooby, E., Tan, K., Zhou, L., King, A., Ramanathan, A., Malhotra, A., Harandi, M., & Marzbanrad, F. (2024). NeoSSNet: Real-time neonatal chest sound separation using deep learning. *IEEE Open Journal of Engineering in Medicine and Biology*, 5, 345–352. <https://doi.org/10.1109/OJEMB.2024.3401571>
- Sun, W., Zhang, Y., & Chen, F. (2024). Research on heart and lung sound separation method based on DAE–NMF–VMD. *EURASIP Journal on Advances in Signal Processing*, 2024, Article 59. <https://doi.org/10.1186/s13634-024-01152-0>
- Zhao, Q., Geng, S., Wang, B., Sun, Y., Nie, W., Bai, B., Yu, C., Zhang, F., Tang, G., Zhang, D., Zhou, Y., Liu, J., & Hong, S. (2024). Deep learning in heart sound analysis: From techniques to clinical applications. *Health Data Science*, 4, Article 0182. <https://doi.org/10.34133/hds.0182>