

CSE6234 Software Design

Software Design Proposal and Software Design Pattern Document

HealthTech

Development of a Machine Learning-Based System for Cardiopulmonary Sound Separation

Group 1

Tutorial Section: TT5L

Name	Student ID
AL-SALOUL, ASHRAF ALI HUSSEIN	1221303805
AHMAD AKMAL ASYRAAF BIN SHAMS	242UC244CJ
RESHMA A/P KRISHNAMURTHY	243UC247KW
SHARWIN A/L R SIVALINGAM	241UC241C1

May 8, 2026

Contents

Abstract / Executive Summary	iv
1 Introduction	1
1.1 Executive Summary / Abstract	1
1.2 Background / Introduction	1
1.3 Problem Statement	2
1.4 Project Objectives	2
1.4.1 Literature Review	3
1.5 Project and System Scope	5
1.5.1 In Scope	5
1.5.2 Out of Scope	5
2 System Overview	6
2.1 Generic Use Case	7
2.2 Class Diagram	7
2.3 Entity Relationship Diagram (ERD)	10
2.4 Object Diagram	11
3 Requirements	13
3.1 Functional Requirements	13
3.2 Non-Functional Requirements	14

3.3	Software Design Concepts	15
3.3.1	Abstraction	15
3.3.2	Modularity	16
3.3.3	Separation of Concerns	16
3.3.4	Information Hiding	16
3.3.5	Low Coupling and High Cohesion	16
3.4	Design Principles	17
3.4.1	Single Responsibility Principle	17
3.4.2	Open/Closed Principle	17
3.4.3	Dependency Inversion Principle	17
3.4.4	Interface-Based Design for Algorithm Replacement	17
3.4.5	Persistence Separation Between SQLite and Filesystem	17
3.5	Software Design Plan	18
3.6	Support from Selected Design Patterns	19
3.6.1	Facade Pattern	19
3.6.2	Strategy Pattern	19
3.6.3	Factory Method Pattern	19
4	Proposed Design Patterns	20
4.1	Facade Pattern	20
4.1.1	Pattern Type	20
4.1.2	Problem	21
4.1.3	Solution	21
4.1.4	Structure	21
4.1.5	Project Application	22
4.2	Strategy Pattern	25
4.2.1	Pattern Type	25

4.2.2	Problem	25
4.2.3	Solution	26
4.2.4	Structure	26
4.2.5	Project Application	27
4.3	Factory Method Pattern	29
4.3.1	Pattern Type	30
4.3.2	Problem	30
4.3.3	Solution	30
4.3.4	Structure	30
4.3.5	Project Application	31
5	Conclusion and Suggestions	35
5.1	Conclusion	35
5.2	Suggestions / Future Work	36
	Bibliography / References	37

Abstract / Executive Summary

This report presents a software design proposal and design pattern document for a HealthTech cardiopulmonary sound separation prototype. The proposed system accepts mixed WAV audio, uses NeoSSNet to separate the recording into heart sound and lung sound outputs, stores structured metadata and processing records in SQLite, and keeps uploaded and generated audio files in local storage.

The report focuses on software design rather than clinical validation. It documents the problem statement, objectives, system scope, UML models, requirements, design principles, and selected design patterns. FastAPI is used as the backend API layer, SQLite supports jobs, results, logs, and metadata, and the local filesystem stores audio artifacts and model files. The Facade, Strategy, and Factory Method patterns are proposed to simplify the separation workflow, support future algorithm extension, and centralise model or algorithm creation.

Chapter 1

Introduction

1.1 Executive Summary / Abstract

This report presents a software design proposal and design pattern document for a cardiopulmonary sound separation system. The proposed system is a HealthTech prototype that accepts mixed WAV audio and separates it into heart sound and lung sound outputs using NeoSS-Net. The project is designed as a local web-based application with a FastAPI backend, a simple HTML/CSS/JavaScript interface, SQLite database storage for metadata and processing records, and local filesystem storage for uploaded and generated audio files.

The purpose of the system is to demonstrate how machine learning inference can be integrated into a reusable software system rather than being kept only as an isolated model script. The application supports upload, validation, real separation, output preview, download, and processing history. This document focuses on software design quality, including requirements, UML modelling, database structure, layered architecture, and design patterns. The prototype is not a clinical diagnosis system and does not claim clinical validation. Instead, it is intended to show a practical and explainable software design for cardiopulmonary sound separation.

1.2 Background / Introduction

Cardiopulmonary sound recordings often contain both heart and lung components because the two sound sources are captured from nearby areas of the chest. In a mixed recording, the sound of one organ can interfere with analysis of the other. This creates a useful software engineering problem: the system must accept an input audio file, apply a sound separation model, produce separate heart and lung audio outputs, and keep the processing records traceable for review.

Recent research has shown that computational methods can support heart and lung sound processing, including source separation and deep learning-based heart sound analysis. However, a model alone is not enough for a usable application. A student, lecturer, researcher, or demonstration user needs a complete workflow that includes file upload, validation, inference, storage, result preview, download, error handling, and history. For this reason, the project is framed as a software design assignment that connects machine learning with maintainable backend and frontend architecture.

The proposed system uses NeoSSNet as the core separation model and stores model-related files in local storage. SQLite is used to store structured data such as uploaded audio metadata, model records, separation jobs, result paths, metrics, and logs. The actual WAV audio files are stored on the filesystem so that the database remains lightweight and easy to inspect. This separation between database records and audio artifacts supports a practical local deployment suitable for a Final Year Project prototype.

1.3 Problem Statement

Mixed cardiopulmonary audio is difficult to use directly when the target task requires separate heart and lung sound outputs. If a system only stores uploaded mixed audio, users cannot easily preview the isolated components or use them for later analysis and demonstration. At the same time, if separation is implemented only as a standalone script, the workflow may lack upload handling, database tracking, result retrieval, logging, and a clear user interface.

The project therefore needs a reusable software system that can connect audio upload, NeoSSNet inference, file storage, SQLite metadata, and browser-based result viewing. The design must remain simple enough for a student project, but structured enough to explain through software design concepts, UML diagrams, and design patterns. The system must not present separated audio as clinical diagnosis; it should be described as a prototype that demonstrates automated cardiopulmonary sound separation and supporting software architecture.

1.4 Project Objectives

The objectives of this project are:

1. To design a local web-based system that accepts mixed cardiopulmonary WAV audio through a browser interface.
2. To validate uploaded audio files before they are stored or processed by the backend.

3. To integrate NeoSSNet as the real cardiopulmonary sound separation model for generating heart and lung output audio.
4. To save separated heart and lung WAV files in local output folders for preview and download.
5. To store structured metadata, job records, result paths, logs, and optional evaluation metrics in SQLite.
6. To provide result viewing, audio preview, download support, and processing history through the web interface.
7. To document a maintainable software design using UML diagrams, requirements, design principles, and selected design patterns.

1.4.1 Literature Review

Cardiopulmonary sound analysis studies the information contained in heart and lung recordings, where useful physiological sounds may overlap because both sources are captured from the chest area. In practical recordings, the heart component can interfere with lung sound analysis, while lung sounds and breathing noise can also obscure heart sound patterns. Sound separation is therefore an important preprocessing and application problem because it attempts to recover clearer heart and lung signals from a mixed recording. For the proposed HealthTech prototype, the literature is relevant not only because it motivates the use of machine learning for heart and lung sound separation, but also because it shows the need for a usable software system around the algorithm. A practical system must support upload, validation, inference, storage, preview, download, and processing history rather than treating separation as an isolated research script.

Deep learning-based separation is the closest research theme to the proposed system because the prototype uses NeoSSNet as its core model. Poh et al. (2024) proposed NeoSSNet for real-time neonatal chest sound separation, with the aim of separating mixed chest sound recordings into heart and lung components. The method follows an audio source separation approach in which a waveform is encoded into a learned representation, separation masks are estimated, and separated outputs are reconstructed. This is directly aligned with the system goal because the application must produce two usable audio files: one heart output and one lung output. NeoSSNet also supports the project from a software design perspective because its inference process can be placed inside a dedicated machine learning layer and called by the backend service layer. However, the research contribution mainly concerns the model and separation performance. It does not by itself provide a complete application workflow for user upload, database-backed job tracking, browser preview, or maintainable backend organization.

A second theme is hybrid separation using signal processing and decomposition methods. Sun et al. (2024) studied a DAE–NMF–VMD method for separating heart and lung sounds. This approach combines deep autoencoder-based representation learning, nonnegative matrix factorization, and variational mode decomposition. Compared with a pure neural separation model such as NeoSSNet, the DAE–NMF–VMD method shows that cardiopulmonary sound separation can also be designed as a multi-stage processing pipeline that combines learned features with mathematical decomposition and denoising. This is useful for the proposed project because it shows that future versions of the system may need to compare different separation approaches. Different algorithms may require different preprocessing steps, configuration parameters, model files, and output handling. Therefore, the software should avoid embedding one algorithm too deeply into the route or controller logic. A modular design with an interchangeable algorithm interface is better suited for future experimentation while keeping the current prototype focused on NeoSSNet.

A broader theme is the role of deep learning in heart sound analysis. Zhao et al. (2024) reviewed deep learning techniques, datasets, and applications for heart sound analysis, including preprocessing, feature extraction, model development, and potential clinical application areas. This review is important because it places the project within a wider research area where machine learning is increasingly used to process phonocardiogram and related biomedical sound data. At the same time, the review also highlights issues such as dataset limitations, generalisation, and the need for further refinement before broad clinical use. These points are important for this report because the proposed system should not be presented as a clinically validated diagnostic tool. It is more accurately described as a proof-of-concept HealthTech software prototype that demonstrates cardiopulmonary sound separation and the supporting software architecture needed to make the workflow usable and traceable.

Taken together, these studies show that cardiopulmonary sound separation and heart sound analysis are active research areas with several promising algorithmic directions. NeoSSNet supports the feasibility of deep learning-based heart and lung sound separation, DAE–NMF–VMD shows the value of hybrid and compositional approaches, and broader heart sound analysis research explains why careful preprocessing, dataset handling, and model evaluation matter. The main gap for this software design report is that existing work focuses strongly on algorithms, while less attention is given to reusable application architecture for a local prototype. In particular, the literature does not fully address database-backed processing history, structured metadata, user upload workflow, separated audio preview and download, local file storage, logging, and maintainability through design patterns. The proposed system responds to this gap by combining FastAPI for the backend API, SQLite for metadata and job records, local filesystem storage for uploaded and separated WAV files, and NeoSSNet for real separation. This design connects research-level separation models with a practical software workflow that can be demonstrated, inspected, and extended in a student software design assignment.

1.5 Project and System Scope

1.5.1 In Scope

The project scope covers a local prototype that accepts mixed WAV audio, validates the file, stores the uploaded audio in a local folder, runs NeoSSNet-based separation, generates separated heart and lung WAV outputs, stores output files locally, records metadata and job status in SQLite, and provides browser-based preview, download, and processing history. The report also covers software design artefacts such as use case modelling, class modelling, database modelling, requirements, design principles, and design patterns.

The dataset and runtime storage are treated as separate concerns. Dataset files such as HLS-CMDS data are used for development or evaluation purposes, while runtime uploads and generated outputs are stored under the application storage folders. This separation avoids mixing training or testing data with user-uploaded demo files.

1.5.2 Out of Scope

The project does not aim to provide medical diagnosis, clinical decision support, hospital deployment, or validated clinical accuracy. It does not claim that separated audio is suitable for real clinical use without further testing and approval. Authentication, cloud deployment, multi-user role management, and production-scale monitoring are also outside the current scope unless they are added in a later phase.

Training a new model from runtime uploads is also outside the current system scope. Runtime uploads are used for inference and demonstration only. The current design focuses on integrating real NeoSSNet inference into a clean software system that is practical for local execution and clear enough to defend in a software design presentation.

Chapter 2

System Overview

The proposed system is a local web-based prototype for cardiopulmonary sound separation. A user uploads a mixed WAV audio file through the browser interface, and the FastAPI backend validates the upload, stores the original file, creates database records, runs NeoSSNet separation, saves the separated heart and lung audio files, and returns results for preview and download. SQLite is used for structured records such as uploaded audio metadata, model information, separation jobs, result paths, metrics, and logs, while local folders are used for the actual audio files and model checkpoints.

The following diagrams describe the system from different software design viewpoints. The use case diagram presents the system functions from the user's perspective, the class diagram shows the planned software components, the entity relationship diagram shows the database structure, and the object diagram gives an example of one completed separation job.

2.1 Generic Use Case




Figure 2.1: Generic Use Case Diagram

Figure 2.1 shows the main actions available to students, lecturers, examiners, researchers, or other demo users. The central use cases are uploading a mixed WAV file, validating the audio, running NeoSSNet separation, previewing the original and separated outputs, downloading the generated heart and lung files, and viewing processing history.

The diagram also shows that several user-facing actions depend on stored system records. For example, result viewing and history display require the system to save metadata, job status, output paths, and logs. This supports the project requirement that the system should not only run separation, but also keep enough information for demonstration, testing, and later review.

2.2 Class Diagram




Figure 2.2: Overall System Class Diagram

Figure 2.2 presents the planned software classes grouped into frontend, router, service, machine learning, and database access responsibilities. The browser interface calls FastAPI routers, the routers delegate work to service classes, and the service layer coordinates storage, database repositories, logging, preprocessing, and model inference.

This design keeps the upload, separation, storage, result viewing, and history features modular. Upload handling is separated from separation orchestration, and the machine learning layer is separated from database access. This makes the system easier to explain and maintain because each class has a focused responsibility rather than placing all workflow logic inside one route function.

2.3 Entity Relationship Diagram (ERD)




Figure 2.3: Entity Relationship Diagram

Figure 2.3 shows the SQLite database structure used to store structured project data. The `uploaded_audio` table records each uploaded WAV file, the `model` table records available

model information, and the `separation_job` table connects an uploaded file with the model used for processing. The `separation_result` table stores the paths and metadata for generated heart and lung outputs.

The design intentionally stores file paths and metadata in SQLite rather than storing large WAV audio blobs in the database. This supports a practical student project architecture: local storage keeps uploaded and generated audio files, while SQLite supports history, result lookup, job status, metrics, and logs. As a result, the browser can display previous jobs and download links without scanning folders manually.

2.4 Object Diagram




Figure 2.4: Object Diagram for a Completed Separation Job

Figure 2.4 gives a concrete example of how the system data may look after one successful separation. One `UploadedAudio` object is linked to one `SeparationJob`, the job uses an active `Model`, and the completed job produces a `SeparationResult` containing heart and lung

output file paths.

The object diagram helps explain how upload, separation, storage, result viewing, and history are connected at runtime. When a user views history, the system can retrieve the job, original upload metadata, model information, output paths, optional metrics, and log record as related objects. This gives the project a traceable flow from the original mixed input to the separated outputs.

Chapter 3

Requirements

This chapter defines the main functional requirements, non-functional requirements, software design concepts, and design principles for the cardiopulmonary sound separation system. The system is treated as a prototype and proof of concept for a software design assignment and Final Year Project context. It is designed to demonstrate a realistic local web application that can upload mixed cardiopulmonary audio, run NeoSSNet-based separation, store metadata, and present results clearly to students, lecturers, researchers, or demonstration users. The requirements do not claim clinical validation; they focus on software behaviour, maintainability, and demonstration readiness.

3.1 Functional Requirements

Functional requirements describe the main behaviours that the system must provide to the user and to the backend workflow.

Table 3.1: Functional Requirements

ID	Requirement	Description
FR-01	Upload mixed WAV audio	The user shall be able to select and upload a mixed cardiopulmonary WAV file through the web interface.
FR-02	Validate audio file	The backend shall check that the uploaded file is a valid supported audio file before saving it for processing.
FR-03	Run cardiopulmonary sound separation	The backend shall run the separation workflow using NeoSSNet to separate the mixed input into heart and lung sound outputs.
FR-04	Generate heart and lung output files	The system shall save separated heart and lung audio files in the configured output folders.
FR-05	Preview separated audio	The web interface shall allow the user to preview the original audio and the separated heart and lung outputs where browser playback is supported.
FR-06	Download separated audio	The user shall be able to download the generated heart and lung output files for later review or demonstration.
FR-07	View processing history	The system shall show previous uploads, jobs, statuses, and result links so that users can review earlier processing activity.
FR-08	Store logs and metadata	The system shall store upload metadata, job records, result paths, metrics where available, and important processing logs in SQLite.

These functional requirements support the main demonstration flow of the project. A user uploads a mixed sound, the backend validates and processes it, the system stores the generated files, and the browser presents previews, downloads, and history. This keeps the project focused on a working end-to-end software system rather than only a standalone machine learning script.

3.2 Non-Functional Requirements

Non-functional requirements describe the quality attributes expected from the prototype. These qualities are important because the system should be understandable, maintainable, and reliable enough for local testing and presentation.

Table 3.2: Non-Functional Requirements

Quality Attribute	Requirement
Usability	The interface should make the upload, separation, preview, download, and history flow clear for a student or examiner during a demonstration.
Maintainability	The code should be organised into clear frontend, router, service, machine learning, database, and storage responsibilities so that future changes are easier to make.
Reliability	The backend should handle invalid files, failed separation jobs, missing outputs, and database errors with clear status messages and logs.
Modularity	Each major part of the system should be replaceable or testable with minimal impact on unrelated components.
Performance	The system should process short demonstration WAV files within a reasonable time on a local machine, while avoiding unnecessary database storage of large audio blobs.
Data integrity	SQLite records should consistently link uploaded audio, model records, separation jobs, results, metrics, and logs. Failed jobs should keep useful error information.
Portability and local execution	The system should run locally using Python, FastAPI, SQLite, PyTorch, and local folders without requiring cloud services or complex deployment infrastructure.

The non-functional requirements guide implementation decisions. For example, local execution supports a practical student project setup, while data integrity and logging support traceability when explaining how a mixed input becomes separated heart and lung outputs.

3.3 Software Design Concepts

3.3.1 Abstraction

Abstraction is used to hide detailed processing steps behind simpler operations. For example, the router should not need to know every audio preprocessing, NeoSSNet inference, file saving, and database update step. It should call a service-level operation that represents the separation workflow at a higher level.

3.3.2 Modularity

Modularity divides the system into smaller parts with clear responsibilities. The frontend handles user interaction, FastAPI routers handle HTTP requests, services coordinate workflows, the ML layer performs separation, the database layer manages SQLite records, and the storage layer manages local files. This structure makes the project easier to understand and easier to extend.

3.3.3 Separation of Concerns

Separation of concerns prevents unrelated responsibilities from being mixed together. Runtime upload folders, generated output folders, model checkpoints, dataset folders, and SQLite meta-data serve different purposes and should remain separate. This design also avoids training from runtime upload folders and avoids storing large WAV content directly in SQLite.

3.3.4 Information Hiding

Information hiding means that internal implementation details are kept inside the module or class responsible for them. File validation rules should remain inside upload or storage-related code, model loading details should remain inside ML-related code, and SQL details should remain inside database access code. Other parts of the system should use clear methods rather than depending on internal details.

3.3.5 Low Coupling and High Cohesion

Low coupling means that components should depend on each other only where necessary. High cohesion means that each component should focus on related tasks. In this project, the service layer may coordinate upload metadata, inference, storage, and result recording, but each detailed operation should stay in the component that owns it. This supports testing and reduces the risk that a small change in one area breaks another area.

3.4 Design Principles

3.4.1 Single Responsibility Principle

The Single Responsibility Principle supports the assignment of one main reason for change to each component. For example, an upload router should change when request handling changes, a storage service should change when file handling rules change, and a NeoSSNet inference class should change when model inference logic changes. This keeps the software design easier to maintain.

3.4.2 Open/Closed Principle

The Open/Closed Principle suggests that the system should be open for extension but closed for unnecessary modification. The separation workflow should be able to support future algorithm options without rewriting the router layer. This is useful if the prototype later compares NeoSSNet with another baseline model or an ONNX version.

3.4.3 Dependency Inversion Principle

Dependency Inversion is relevant because high-level workflow classes should not be tightly bound to one concrete model implementation. A service can depend on an abstract algorithm interface, while the NeoSSNet strategy provides the current concrete implementation. This keeps the workflow stable even if the underlying model class changes.

3.4.4 Interface-Based Design for Algorithm Replacement

Interface-based design supports future algorithm replacement by defining the expected input and output of a separation algorithm. The current algorithm should produce separated heart and lung audio outputs, but the service layer should not need to know whether the implementation uses NeoSSNet, an ONNX-exported model, or a future baseline method. For the current project, NeoSSNet remains the required real separation model.

3.4.5 Persistence Separation Between SQLite and Filesystem

The design separates structured metadata from large binary audio files. SQLite stores records such as filenames, file paths, upload timestamps, job status, model information, output paths,

metrics, and logs. The filesystem stores uploaded WAV files, generated heart and lung outputs, model checkpoints, and dataset files. This approach is simpler to run locally and avoids making the database unnecessarily large.

3.5 Software Design Plan

The proposed software design is organised into layers. This layered plan supports a clear implementation path and makes the system easier to explain in the report and viva presentation.

Table 3.3: Software Design Plan by Layer

Layer	Main Responsibility	Project Role
Frontend layer	HTML, CSS, JavaScript, templates, and browser audio controls	Provides upload form, status messages, original and separated audio playback, download links, and history display.
Backend API layer	FastAPI application and routers	Receives upload, separation, result, download, history, log, and health requests.
Service layer	Workflow coordination and business logic	Connects validation, database updates, NeoSSNet inference, output storage, logging, and result retrieval.
ML inference layer	Audio preprocessing and model inference	Loads or prepares audio, applies the NeoSSNet model, and returns separated heart and lung waveforms.
Database layer	SQLite schema and repository-style access	Stores uploaded audio metadata, model records, job states, result paths, metrics, and system logs.
Storage layer	Local filesystem folders	Stores uploaded mixed WAV files, temporary files, separated heart and lung outputs, logs where needed, and ML model weights.

The design plan follows the project priority of real functionality first, then simplicity and local execution. It avoids unnecessary cloud, microservice, or authentication complexity unless those features are explicitly added later.

3.6 Support from Selected Design Patterns

3.6.1 Facade Pattern

The Facade Pattern supports subsystem simplification. The separation workflow includes validation, preprocessing, NeoSSNet inference, file saving, database updates, and logging. A facade such as `SeparationService` can present one clear operation to the router while coordinating these internal subsystem steps.

3.6.2 Strategy Pattern

The Strategy Pattern supports future algorithm flexibility. Although NeoSSNet is the core model for this project, a strategy interface allows the system design to describe how another separation algorithm could be introduced later without changing the route-level workflow. This is useful for experimentation while still keeping NeoSSNet as the main real separation method.

3.6.3 Factory Method Pattern

The Factory Method Pattern supports model or algorithm object creation. A factory can create the correct `SeparationAlgorithm` based on active model configuration or database metadata. This reduces hardcoding and keeps model construction logic in one place, which is easier to maintain if the project later adds another model version or deployment format.

Chapter 4

Proposed Design Patterns

This chapter proposes three design patterns for the cardiopulmonary sound separation system. The patterns are presented as practical backend design structures for the FastAPI, SQLite, local storage, and NeoSSNet pipeline. Each pattern is linked to a specific project problem so that the design remains useful for implementation and explanation during the final presentation.

The class and method names used in this chapter are aligned with the current model-selection implementation. The report uses `SeparationService` to describe the service role implemented in `app/services/separation_service.py`, while the actual code also includes `SeparationAlgorithm`, `NeoSSNetStrategy`, `SeparationAlgorithmFactory`, and `ModelService`.

4.1 Facade Pattern

The Facade Pattern is proposed to provide a simple interface over the separation workflow, which contains several lower-level operations such as audio preprocessing, model inference, file storage, database updates, and logging (Gamma et al., 1994).

4.1.1 Pattern Type

The Facade Pattern is a structural design pattern. It focuses on arranging classes so that a client can access a complex subsystem through a simpler and more stable interface.

4.1.2 Problem

The separation endpoint should not directly coordinate every step of the backend workflow. If the FastAPI router directly calls the audio preprocessor, NeoSSNet inference class, storage utilities, database session, and logging logic, the route becomes difficult to read, test, and modify. This would also make the design harder to explain because web routing logic and machine learning workflow logic would be mixed in one place.

4.1.3 Solution

The solution is to use `SeparationService` as the facade for the separation workflow. The router calls one high-level operation, `separate_uploaded_audio(db, audio_id, model_id)`, while the service coordinates model selection, algorithm creation, job creation, inference, output storage, and database result updates. This keeps the router focused on request handling and keeps the separation workflow inside a dedicated service layer.

4.1.4 Structure

The template structure below shows how a facade sits between a client and a group of subsystem classes. The client depends on the facade instead of depending on each subsystem class directly.




Figure 4.1: Facade Pattern Template Structure

1. **Client:** The external caller that needs to use the subsystem. In this project, the client is mainly the FastAPI router that receives a separation request.

2. **Facade:** The simplified interface used by the client. In this project, `SeparationService` hides the detailed separation steps behind one service method.
3. **Additional Facade:** An optional second facade can be used when the subsystem grows. The current implementation keeps model lookup in `ModelService`, so no additional facade is required at this stage.
4. **Complex Subsystem:** The lower-level components that perform the actual work. In this project, these include `ModelService`, `SeparationAlgorithmFactory`, `SeparationAlgorithm`, `NeoSSNetStrategy`, SQLite database records, and local file storage.

4.1.5 Project Application

i. Function or Software Component Affected

The affected component is the backend separation workflow, especially the service layer that connects the FastAPI separation route to model selection, machine learning inference, SQLite records, and local file storage. The main facade role is `SeparationService`, represented by the `separate_uploaded_audio` workflow in `app/services/separation_service.py`.

ii. Description of Workflow or Data Flow

The router receives a request to separate an uploaded WAV file and may also receive an optional `model_id`. It passes the database session, uploaded audio identifier, and selected model identifier to `SeparationService`. The service loads the upload metadata from SQLite, selects either the requested model or the active model through `ModelService`, creates the correct algorithm through `SeparationAlgorithmFactory`, runs the selected `SeparationAlgorithm`, saves separated heart and lung WAV files to local storage, and records the job and result metadata in SQLite.

iii. Sample Class Diagram




Figure 4.2: Facade Pattern Project Class Diagram

iv. Sample Other UML Notation




Figure 4.3: Facade Pattern Sequence Diagram

v. Sample Potential Code

```
# Client: FastAPI route in app/routers/separation.py
def separate_audio(audio_id: int, model_id: int | None, db):
    return SeparationService().separate_uploaded_audio(
        db=db,
        audio_id=audio_id,
        model_id=model_id,
    )

# Facade: app/services/separation_service.py
class SeparationService:
    def separate_uploaded_audio(self, db, audio_id, model_id=None):
        upload = get_uploaded_audio(db, audio_id)

        # Subsystem: model selection
        model = ModelService.get_model_for_separation(db, model_id)

        # Subsystem: factory creates the selected ML algorithm
        algorithm = SeparationAlgorithmFactory.create(model)

        # Subsystem: database job and local file paths
        job = create_running_job(
            db,
            upload.uploaded_audio_id,
            model.model_id,
        )
        heart_path = make_heart_output_path(job.job_id)
        lung_path = make_lung_output_path(job.job_id)

        # Subsystem: audio/ML algorithm
        result = algorithm.separate(
            input_wav_path=resolve_project_path(upload.stored_path),
            model_path=resolve_project_path(model.checkpoint_path),
            model_config_path=resolve_project_path(model.config_path),
            heart_output_path=heart_path,
            lung_output_path=lung_path,
        )
```

```
save_result_metadata(db, job, result)
return job.job_id
```

The code sample is simplified for readability, but the roles and method names follow the implemented backend structure.

vi. Benefits

The facade makes the route simpler, improves separation of concerns, and makes the workflow easier to test as one service operation. It also improves presentation clarity because the main separation process can be explained as a single service call that coordinates several internal components.

vii. Limitations

The main limitation is that the facade can become too large if every processing rule is placed inside `SeparationService`. The implementation should keep detailed work inside subsystem classes and use the facade mainly for orchestration.

4.2 Strategy Pattern

The Strategy Pattern is used to keep the separation algorithm interchangeable behind a common interface. This is useful because the project currently uses NeoSSNet, but the backend should still allow another implemented algorithm to be added later without changing the upload, result, history, or download routes (Gamma et al., 1994).

4.2.1 Pattern Type

The Strategy Pattern is a behavioral design pattern. It focuses on defining a family of algorithms and allowing the context object to use them through a shared interface.

4.2.2 Problem

The system should not hardcode NeoSSNet-specific logic throughout the service layer. If model-specific preprocessing, inference, and output handling are tightly coupled to the route or

service, future experiments become harder. Even for a student project, keeping the algorithm boundary clear makes the design easier to test and explain.

4.2.3 Solution

The solution is to define a SeparationAlgorithm interface with a separate(...) operation. NeoSSNetStrategy is the current concrete strategy and wraps the existing real run_neossnet_inference function. SeparationService acts as the context and calls the selected strategy through the shared interface.

4.2.4 Structure

The template structure below shows a context object that delegates algorithm-specific work to a strategy object. The client uses the context without needing to know which concrete strategy performs the work.




Figure 4.4: Strategy Pattern Template Structure

1. **Context:** The object that uses a strategy to complete part of its work. In this project, SeparationService is the context because it delegates model-specific separation to an algorithm object.
2. **Strategy Interface:** The common contract for all algorithms. In this project, SeparationAlgorithm defines the expected separation method and output format.
3. **Concrete Strategies:** The actual algorithm implementations. In the current implementation, NeoSSNetStrategy is the concrete strategy.

4. **Client:** The caller that starts the workflow. In this project, the FastAPI router starts the separation request, while the service uses the configured strategy internally.

4.2.5 Project Application

i. Function or Software Component Affected

The affected component is the machine learning inference boundary. The main classes are `SeparationAlgorithm`, `NeoSSNetStrategy`, and `SeparationService`. The interface is intentionally small so that a new implemented strategy can be registered later without changing the separation route or result workflow.

ii. Description of Workflow or Data Flow

The service resolves the uploaded WAV path and model paths, then calls the selected `SeparationAlgorithm`. The current concrete strategy, `NeoSSNetStrategy`, delegates to the existing NeoSSNet inference function and returns a `SeparationAlgorithmResult` containing heart output path, lung output path, sample rate, duration, and file size metadata. The service records these outputs without depending on NeoSSNet internals.

iii. Sample Class Diagram




Figure 4.5: Strategy Pattern Project Class Diagram

iv. Sample Other UML Notation




Figure 4.6: Strategy Pattern Sequence Diagram

v. Sample Potential Code

Strategy: app/ml/separation_algorithm.py

```
class SeparationAlgorithm(Protocol):
```

```
    def separate(
        self,
        input_wav_path,
        model_path,
        model_config_path,
        heart_output_path,
        lung_output_path,
        device_name="cpu",
    ) -> SeparationAlgorithmResult:
        ...
```

ConcreteStrategy: app/ml/neossnet_strategy.py

```
class NeoSSNetStrategy:
```

```
    def separate(self, input_wav_path, model_path, model_config_path,
                 heart_output_path, lung_output_path, device_name="cpu"):
        return run_neossnet_inference(
            input_wav_path=input_wav_path,
            model_path=model_path,
            model_config_path=model_config_path,
```

```

        heart_output_path=heart_output_path,
        lung_output_path=lung_output_path,
        device_name=device_name,
    )

# Context: app/services/separation_service.py
class SeparationService:
    def run_selected_algorithm(self, algorithm: SeparationAlgorithm):
        return algorithm.separate(
            input_wav_path=input_path,
            model_path=model_path,
            model_config_path=model_config_path,
            heart_output_path=heart_output_path,
            lung_output_path=lung_output_path,
        )

```

The concrete strategy does not duplicate NeoSSNet inference logic; it delegates to the existing inference function.

vi. Benefits

The strategy design makes the algorithm boundary clear, improves testability, and supports future comparison with alternative separation methods. It also helps keep NeoSSNet-specific code inside a dedicated class rather than spreading it across the router and service layer.

vii. Limitations

The limitation is that the project may only use NeoSSNet during the final submission. If no alternative algorithm is implemented, the extra interface should remain small so it does not make the student project unnecessarily complicated.

4.3 Factory Method Pattern

The Factory Method Pattern is proposed to centralise the creation of separation algorithm objects. This avoids scattering model creation, checkpoint loading, and strategy selection logic throughout the backend (Gamma et al., 1994).

4.3.1 Pattern Type

The Factory Method Pattern is a creational design pattern. It focuses on creating objects through a dedicated creator method instead of directly constructing concrete classes in client code.

4.3.2 Problem

The system stores model-related information such as model name, version, checkpoint path, and active status in SQLite or local configuration. If each service directly constructs `NeoSSNetStrategy` and loads checkpoint paths manually, model creation becomes duplicated and harder to change. This is especially risky if the project later supports more than one model format.

4.3.3 Solution

The solution is to use `SeparationAlgorithmFactory` as the creator. The service obtains the selected or active model through `ModelService`, then passes that model record to the factory. The factory checks the model architecture and returns a `SeparationAlgorithm` object. At the current stage, the concrete product is `NeoSSNetStrategy`.

4.3.4 Structure

The template structure below shows a creator that returns a product interface while hiding the exact concrete product class from the client.




Figure 4.7: Factory Method Pattern Template Structure

1. **Creator:** The class or method responsible for creating product objects. In this project, SeparationAlgorithmFactory is the proposed creator.
2. **Concrete Creator:** The specific creation logic that knows which concrete product to instantiate. In this project, the factory method can create a NeoSSNet-based strategy from the active model configuration.
3. **Product:** The shared interface returned to the client. In this project, SeparationAlgorithm is the product interface.
4. **Concrete Product:** The actual object created by the factory. In this project, NeoSSNetStrategy is the main concrete product.
5. **Client:** The class that requests a product from the factory. In this project, SeparationService uses the factory so that it does not need to hardcode model construction details.

4.3.5 Project Application

i. Function or Software Component Affected

The affected component is model and algorithm construction. The implemented creator is SeparationAlgorithmFactory, and it is used by SeparationService. Model lookup is handled separately by ModelService.

ii. Description of Workflow or Data Flow

The service requests the selected or active model from `ModelService`. It then passes the model record to `SeparationAlgorithmFactory.create(model)`. The factory checks `model.architecture`; when it is `NeoSSNet`, it creates `NeoSSNetStrategy`. The created object is returned through the `SeparationAlgorithm` interface, so the service can call `separate(...)` without hardcoding the concrete strategy construction.

iii. Sample Class Diagram




Figure 4.8: Factory Method Pattern Project Class Diagram

iv. Sample Other UML Notation




Figure 4.9: Factory Method Pattern Sequence Diagram

v. Sample Potential Code

```

# Product: app/ml/separation_algorithm.py
class SeparationAlgorithm(Protocol):
    def separate(self, input_wav_path, model_path, model_config_path,
                  heart_output_path, lung_output_path, device_name="cpu"):
        ...

# ConcreteProduct: app/ml/neossnet_strategy.py
class NeoSSNetStrategy:
    def separate(self, input_wav_path, model_path, model_config_path,
                  heart_output_path, lung_output_path, device_name="cpu"):
        return run_neossnet_inference(...)

# Creator: app/services/model_factory.py
class SeparationAlgorithmFactory:

```

```

_registry = {
    "neossnet": NeoSSNetStrategy,
}

@classmethod
def create(cls, model) -> SeparationAlgorithm:
    architecture = model.architecture.strip().lower()
    strategy_class = cls._registry.get(architecture)
    if strategy_class is None:
        raise UnsupportedModelArchitectureError(architecture)

    return strategy_class()

# Client: app/services/separation_service.py
class SeparationService:
    def get_algorithm(self, db, model_id=None):
        model = ModelService.get_model_for_separation(db, model_id)
        return SeparationAlgorithmFactory.create(model)

```

The factory creates the strategy object only. The model checkpoint and configuration paths remain in the model record and are passed into the strategy during the separation call.

vi. Benefits

The factory method reduces hardcoded model creation, keeps checkpoint loading in one place, and makes it easier to add a future model type without changing every service that needs an algorithm. It also supports cleaner testing because a test can provide a simple model configuration and check which strategy is created.

vii. Limitations

The limitation is that a factory may be unnecessary if the system permanently supports only one model. For this project, the factory should stay small and practical, mainly to keep model construction clear and centralised.

Chapter 5

Conclusion and Suggestions

5.1 Conclusion

This report proposed the design of a HealthTech cardiopulmonary sound separation system as a prototype and proof of concept. The system is intended to accept mixed WAV audio, run cardiopulmonary sound separation, and produce separate heart sound and lung sound output files. It supports the main user workflow of upload, validation, separation, preview, download, and processing history while keeping the implementation practical for a student software design project.

The proposed architecture is important because the machine learning model is only one part of the complete system. FastAPI provides the backend API layer for handling upload, separation, result, download, and history requests. NeoSSNet provides the core separation model for generating heart and lung audio outputs. SQLite stores structured metadata such as uploaded audio records, model details, separation jobs, result paths, evaluation metrics, and logs. Local storage keeps the actual uploaded WAV files, separated output files, and model weights, which avoids storing large audio content directly inside the database.

The selected design patterns strengthen the proposed design. The Facade Pattern simplifies the separation workflow by allowing the router to call a high-level service instead of directly managing preprocessing, inference, storage, logging, and database updates. The Strategy Pattern supports future algorithm flexibility by allowing the separation model to be represented through a common interface. The Factory Method Pattern centralises the creation of model or algorithm objects so that model selection and construction are not scattered across the codebase.

Overall, the proposed design demonstrates how machine learning, database persistence, local file storage, and web interaction can be organised into a reusable software system. The report does not claim clinical validation or real hospital deployment. Instead, it presents a clear and

maintainable software design suitable for assignment submission, demonstration, and future implementation improvement.

5.2 Suggestions / Future Work

Several improvements can be considered after the core prototype is working:

1. **Improve frontend usability:** The user interface can be refined with clearer progress indicators, better error messages, improved history filtering, and more polished audio playback controls.
2. **Add more datasets and noise conditions:** Future testing can include broader cardiopulmonary recordings, different noise levels, and additional real-world recording conditions to better understand model behaviour.
3. **Add asynchronous or background processing:** Longer audio files or slower inference tasks can be handled using background jobs so that the web request does not need to wait synchronously for the full separation process.
4. **Add authentication if needed:** If the system is extended beyond a local demonstration, user login and role-based access can be added to protect uploaded files and processing history.
5. **Add algorithms through Strategy:** Alternative approaches, baseline methods, or future ONNX-based versions can be added as new strategies without changing the main router workflow.
6. **Improve evaluation metrics and reporting:** The system can provide clearer metric reports such as SNR, SDR, SI-SDR, MSE, or MAE where suitable reference data is available.
7. **Validate on broader real-world recordings:** Further validation should be performed on more varied recordings before making any claim about robustness or clinical usefulness.

In conclusion, the project provides a strong foundation for a practical software design around cardiopulmonary sound separation. Its main value is not only the use of NeoSSNet, but also the structured design that makes upload handling, inference, storage, logging, result viewing, and future extension easier to understand and maintain.

Bibliography / References

- Gamma, E., Helm, R., Johnson, R., & Vlissides, J. (1994). *Design patterns: Elements of reusable object-oriented software*. Addison-Wesley.
- Poh, Y. Y., Grooby, E., Tan, K., Zhou, L., King, A., Ramanathan, A., Malhotra, A., Harandi, M., & Marzbanrad, F. (2024). NeoSSNet: Real-time neonatal chest sound separation using deep learning. *IEEE Open Journal of Engineering in Medicine and Biology*, 5, 345–352. <https://doi.org/10.1109/OJEMB.2024.3401571>
- Sun, W., Zhang, Y., & Chen, F. (2024). Research on heart and lung sound separation method based on DAE–NMF–VMD. *EURASIP Journal on Advances in Signal Processing*, 2024, Article 59. <https://doi.org/10.1186/s13634-024-01152-0>
- Zhao, Q., Geng, S., Wang, B., Sun, Y., Nie, W., Bai, B., Yu, C., Zhang, F., Tang, G., Zhang, D., Zhou, Y., Liu, J., & Hong, S. (2024). Deep learning in heart sound analysis: From techniques to clinical applications. *Health Data Science*, 4, Article 0182. <https://doi.org/10.34133/hds.0182>